

Preliminary Data Access Center : User Report

Krzysztof Suberlak,¹★ Željko Ivezić,¹ and the PDAC team

¹*Department of Astronomy, University of Washington, Seattle, WA, United States*

Accepted XXX. Received YYY; in original form ZZZ

ABSTRACT

A report on user experience of the Preliminary Data Access Center (PDAC). We test the quality and ease of access to the data. PDAC will pave the way to the Science User Interface and Tools (SUIT). We employ both in-detail study of individual objects, and a statistical study of an ensemble of objects. We evaluate user-friendliness of the current interface, and make recommendations for its future improvements.

1 INTRODUCTION

This is a document to report on the user experience testing of the Preliminary Data Access Center. The Large Scale Synoptic Telescope (LSST) will produce a big volume of data. Such unprecedented data stream poses new challenges to provide an easy access for users, in such a way that they can quickly find what they need, and thus be able to focus on the science goal that they would like to achieve. The detail description of such online user-interface called Science User Interface and Tools is outlined in documents LDM-130 (SUIT requirements) and LDM-492 (SUIT Vision). An idea of having an interface to the data is not new : there exists Aladin, SDSS CAS jobs, IPAC IRSA, Mikulsky NASA Archive, NED, and many other archives. These allow a user to query for data (either via SQL query, or interface), returning the data table. Some user interfaces (eg. IRSA) have some rudimentary plotting capabilities. There have been ideas of a new interface, that would not only eg. plot the lightcurve and display the spectrum, but also allow the user to run some machine learning algorithms, or simple models that can help narrow down the query, or obtain science results in the browser. Namely, Victor Pankratius, from MIT, in his talk "Computer-Aided Discovery: Towards Scientific Insight Generation with Machine Support" outlined the idea of an ipython notebook - access to data, which lives in the cloud, is allocated some CPU share and memory, and allows one to upload / download the data and run the model in real time, which is especially helpful to geoscientists doing fieldwork, where new data acquisition conditions their next step.

These requirements and the vision for SUIT have been further described on confluence pages¹. Some technical notes about current implementation of SUIT by PDAC are also available via confluence pages².

2 OVERVIEW OF PERFORMED TESTS

We test a variety of aspects of PDAC : the user interface, infrastructure, and database ingestion. The user interface is similar to IRSA, which aids the ease of access. In Section 3 we describe the functionality available through user interface. It is a work in progress, hence any deficiency outlined may become updated in real-time, whereas some recommendations, if met with approval, may have a longer implementation timescale. In Section 4 we describe the structure of available data : both data that is available directly from NCSA (internal catalogs) , and data that is available from IRSA (external catalogs). In that section we also provide an overview of query and analysis methods available directly through the User Interface, as well as through SQL. Finally, in Section 5 we consider the quality of database ingestion, answering the question of how well was a given dataset loaded into PDAC. In particular we compare the S82 forced photometry dataset, an outcome of the Summer 2013 reprocessing, to the same data stored locally at the University of Washington.

3 USER INTERFACE: WHAT WE SEE

In order to access PDAC we follow the directions² that include logging to NCSA via VPN <https://vpn.ncsa.illinois.edu/> using Cisco AnyConnect Secure Mobile Client, and opening in the web browser <http://lsst-sui-proxy01.ncsa.illinois.edu/suit>. This opens the main interface screen, which allows to select the database, and perform the desired query.

Currently, PDAC v1, in the upper-left corner of the interface, under tab 'LSST Data' (see Fig. 1) includes the Summer 2013 DM-stack reprocessed SDSS Stripe 82 data (database name `sdss_stripe82_00`) , hosted at the NCSA on the LSST prototype ("integration cluster") hardware, in Qserv [Gregory Dubois-Felsmann, priv.comm. 02-20-2017, slack]. The only other locally stored database (as of March 2017), is WISE catalog, that is not yet accessible via the graphical user interface (it can be queried as Data Base `wise_00`, with catalogs 'Object' containing objects (like

¹ <https://confluence.lsstcorp.org/display/DM/Science+User+Interface+and+Tools>

² <https://confluence.lsstcorp.org/display/DM/Guide+to+PDAC+version+1>

DeepSource in S82 above), and 'ForcedSource' containing forced photometry (like DeepForcedSource in S82)).

The upper-left corner of the interface also leads to 'External Images' and 'External Catalogs'. The Catalogs are all NASA/IPAC³ Infrared Science Archive (IRSA) publicly accessible catalogs, including GAIA, WISE, 2MASS, SPITZER, etc. (see Fig. 2).

4 INFRACTRUCTURE : WHAT IS AVAILABLE AND HOW TO GET IT

As we described in Section 3, the main user interface allows access to the internally stored (at NCSA) SDSS Stripe 82 data reprocessed during the Summer 2013⁴ as part of Data Challenge with the continuously developed LSST Stack⁵.

The reprocessing included:

- coadding the data from all epochs in each of the ugriz SDSS filters. Measurements on coadds (per object) are available as `RunDeepSource` table, accessible via Catalogs → 'DeepSource'. The single-band coadded images with MariaDB metadata are available as `DeepCoadd` table, accessible via Images → 'DeepCoadd'.
- using i-band detections to seed forced photometry on all epochs in all bands. The results of photometry are available as `RunDeepForcedSource` table, accessible via Catalogs → 'Deep Forced Source'.
- For reference, the individual calibrated single epoch images are available as `Science_Ccd_Exposure` table, accessible via Images → 'Science CCD Exposure'.

Additional details of the schema are also outlined in the LSST Data Challenge Report [Shaw, Juric, Becker, Krughoff et al. 2013], and the LSST Database Schema Browser⁶.

Spatial queries that can be directly executed from the PDAC interface, called 'Method Search', include cone, box, elliptical and polygon (See Fig. 1). Spatial queries allow to choose a certain region of the sky by the object ra,dec coordinates. Cone, elliptical, and box queries return objects in a region of the sky bound by a geometrical shape centered on given coordinates (ra,dec). Cone is the most useful type of query, allowing to find objects within a certain radius from the coordinate query. Elliptical search allows to define the shape by an ellipse with a given semi-major axis, position angle and the axis ratio. A box is a square centered on the query coordinates, with a given side size. A polygon allows to define the search region by between 3 and 15 coordinate pairs (vertices of the polygon). Note : Multi-object query is listed in the drop-down menu, but has not yet been implemented (March 2017) - in the future it will allow the user to upload a list of ra,dec and search radii, finding 1-to-1 matches in the existing catalog. An All-Sky option (no spatial constraints) has not been tested given the size of the database.

³ Infrared Processing and Analysis Center, <http://www.ipac.caltech.edu/project/lsst>

⁴ <https://confluence.lsstcorp.org/display/DM/Properties+of+the+2013+SDSS+Stripe+82+reprocessing>

⁵ <https://pipelines.lsst.io/index.html>

⁶ https://lsst-web.ncsa.illinois.edu/schema/index.php?t=DeepForcedSource&sVer=S12_lsstsim

Any query returns a list of all objects within the given region (Fig. 3).

4.1 Postage Stamp Miniatures

We compared the postage stamp miniatures showing the overview of the region against which a given query was performed. We find that the miniature image does not always come from the catalog we query against. In fact, the "coverage" image comes from IRAS, DSS, 2MASS, or WISE - the survey is chosen depending on the size of the region needed to be shown [Xiuqin Wu, priv.comm., 2017]. Indeed, as the query region is increased, the shown image changes unexpectedly from DSS to IRAS or WISE, without issuing a relevant information to the user. A recommendation is to display information about which sky survey a particular image is coming from.

As a concrete example, we execute a cone query against the Deep Source table, expecting to find a point source: a star at ra,dec = 23h30m57.31s, +1d1m13.8s (or 352.73878967464526, 1.020496777615987 degrees), with search radius of 2, 10, 100 and 1000 arcseconds. We find that the queried region only sometimes has the postage stamp miniature aligned with the query region.

We also consider a galaxy located at ra,dec = 40.433, 0.449. A 10 arcsec cone query only sometimes is centered on the galaxy.

4.2 Time Series : periodogram

We compare the g-band time-series of a few objects, taking as the ground truth periodicities reported by Sesar et al. (2010). We record the execution time, and describe in detail the experience of exploring the periodogram tool.

Consider RR Lyrae star id=13350 (ra, dec = 0.283437, 1.178522). A cone query with search radius of 2 arcseconds brings us to a lightcurve view, from where it is possible to open the Time Series view (see Fig.4)

Finally, we compare the PDAC Time Series User Interface to that of the NASA Exoplanet Archive Periodogram⁷ (see Fig. 7). Using few RR Lyrae PDAC g-band lightcurves, each calculation is allocated time slot of approximately 15 seconds. Also see Table 1 for a summary of results.

5 DATABASE INGESTION : IS WHAT WE GET WHAT WE EXPECTED TO GET?

5.1 Positional comparison : box query

We compare the ra,dec of objects found within a given region in PDAC S82 to the analogous data stored at UW (DeepSource tables). Regardless of photometry, for a given objectId, the ra,dec should agree.

⁷ <http://exoplanetarchive.ipac.caltech.edu/cgi-bin/Pgram/nph-pgram>

Figure 1. The main user interface of PDAC ver. 1

Figure 2. IPAC-hosted catalogs, accessible via IRSA.

5.2 Light curve comparison : multiple cone queries

We perform multiple cone queries given known coordinates of well-classified sources. We take as ground truth 484 RR Lyrae stars from [Sesar et al. \(2010\)](#), which were identified by color cuts, Lomb-Scargle periodogram and then confirmed by template fitting. Thus the periods from that work are considered to be reliable (see RR Lyrae 'ground truth' period distribution on Fig.8). Both fit parameters and SDSS lightcurves used in [Sesar et al. \(2010\)](#) are publicly accessible in the online version of the journal.

We perform individual SQL cone queries against PDAC for each RR Lyrae, allowing a generous 2 arcsec search radius (see Appendix A). Location of these RR Lyrae on the sky is shown on Fig. 13.

Lomb-Scargle periodogram does not always find the 'true' period - it is subject to non-uniform sampling, aliasing, and necessity of choosing well the frequency sample on which periodogram powers are evaluated (see [VanderPlas \(2017\)](#) for a recent overview). We nevertheless perform few simple sanity checks :

- (i) Does the PDAC lightcurve folded on 'true' period look real?
- (ii) Using naive Lomb-Scargle, do we find the same period with the S82 Summer 2013 locally-stored data, and PDAC-hosted equivalent dataset?
- (iii) Using naive Lomb-Scargle, do we find the same period with the [Sesar et al. \(2010\)](#) data and PDAC data?

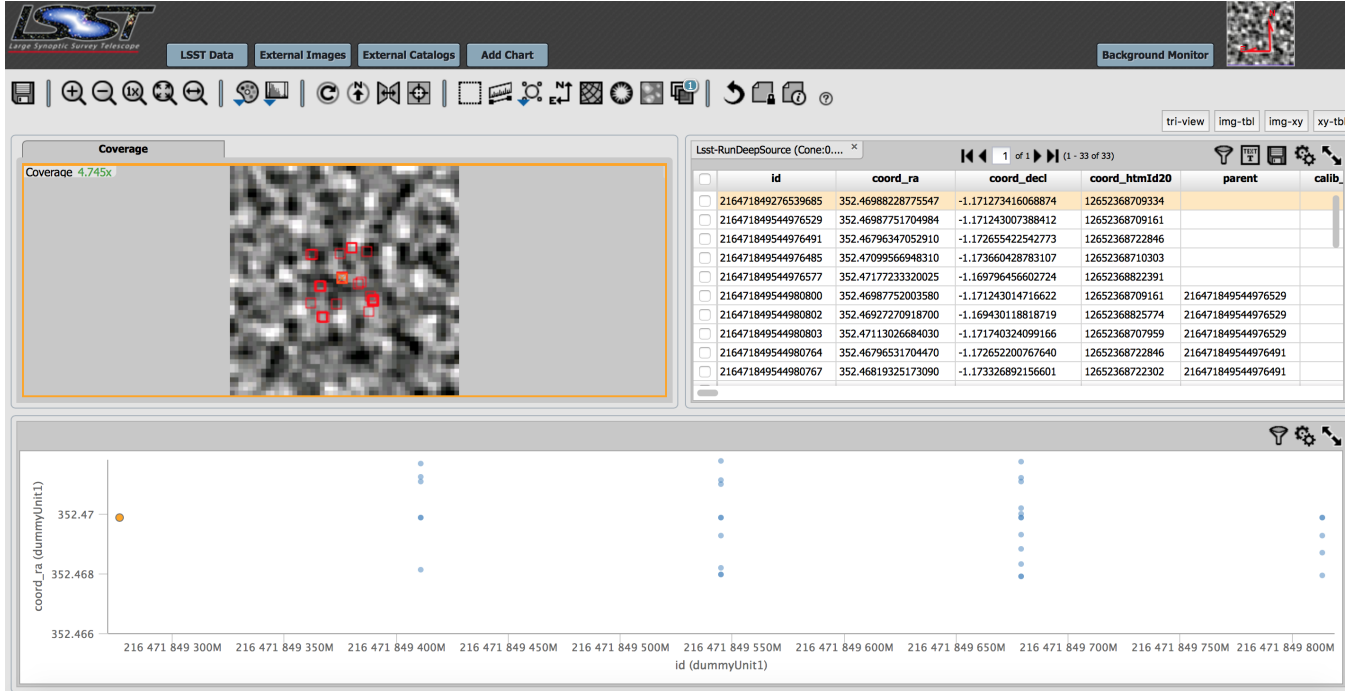


Figure 3. Example cone query against Deep Source table, returning all objects detected within a radius of 10 arcseconds from the position $ra=352.469872$, $dec=-1.171239$. Note that the background image (postage stamp miniature) does not show the actual S82 coadds. This particular feature is described further in Sec. 4.1

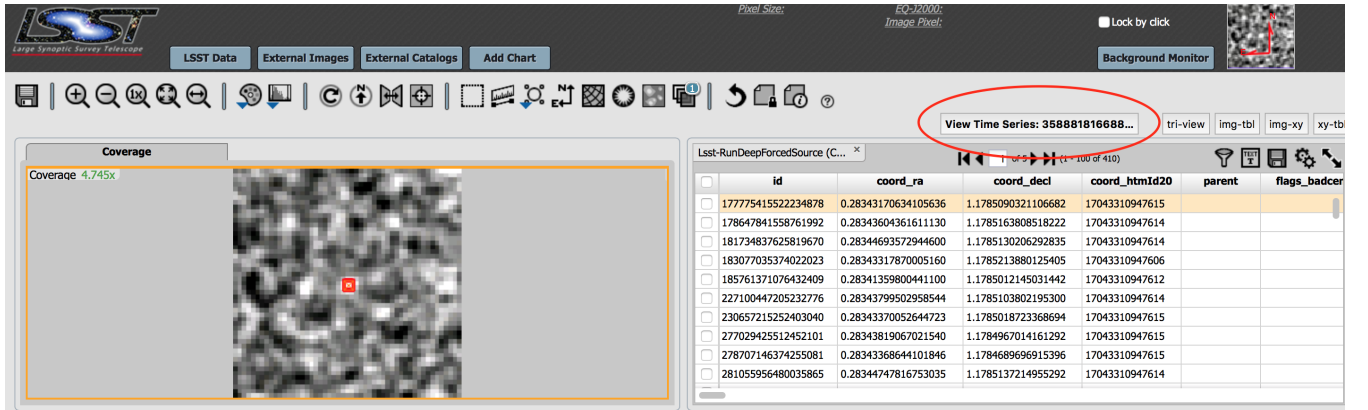


Figure 4. A result of cone query against Run Deep Forced Source table (containing S82 S13 data), with $(ra, dec, radius) = (0.283437, 1.178522, 0.00055)$ degrees). We circle the 'View Time Series' button that links to the Time Series UI shown on Fig. 5.

Question 3 is less direct than questions 1 and 2, since Sesar et al. (2010) used an earlier SDSS Data Release than what was used to create the S82 S13 reprocessed dataset. However, question 1 test solely whether the same object is stored in PDAC as we would assume (since it is unlikely that a random object would be well-represented by an RR Lyrae period). Similarly, Question 2 employs datasets that should be identical, and therefore we would assume that Lomb-Scargle tests would yield identical periodograms.

Assuming that the periods for the 483 RR Lyrae in Sesar et al. (2010) are correct, we attempt to download PDAC data for these objects, compute periodograms, and fold them on the true periods. This serves both to confirm

that the periods are correct, and that the objects themselves are indeed RR Lyrae (see Sec. 5.2)

Using *astroML* python module (Vanderplas et al. 2012), we sample the uniformly spaced frequency grid with $N=5000$ samples span between the smallest and the largest frequency reported in Table 1 of Sesar et al. (2010) $\pm 10\%$, i.e. $\omega_{min} = 0.9(2\pi/P_{min})$, $\omega_{max} = 1.1(2\pi/P_{min})$. We use the default *astroML* Lomb Scargle periodogram settings, namely generalized LS (see Eq.20 in Zechmeister & Kürster (2009), and Section 10.3.2 in Ivezić et al. (2014)).

Using the same frequency grid for all 483 RR Lyrae, we compute Lomb-Scargle periodograms, and determine the best-fit period from the highest frequency peak (see eg.

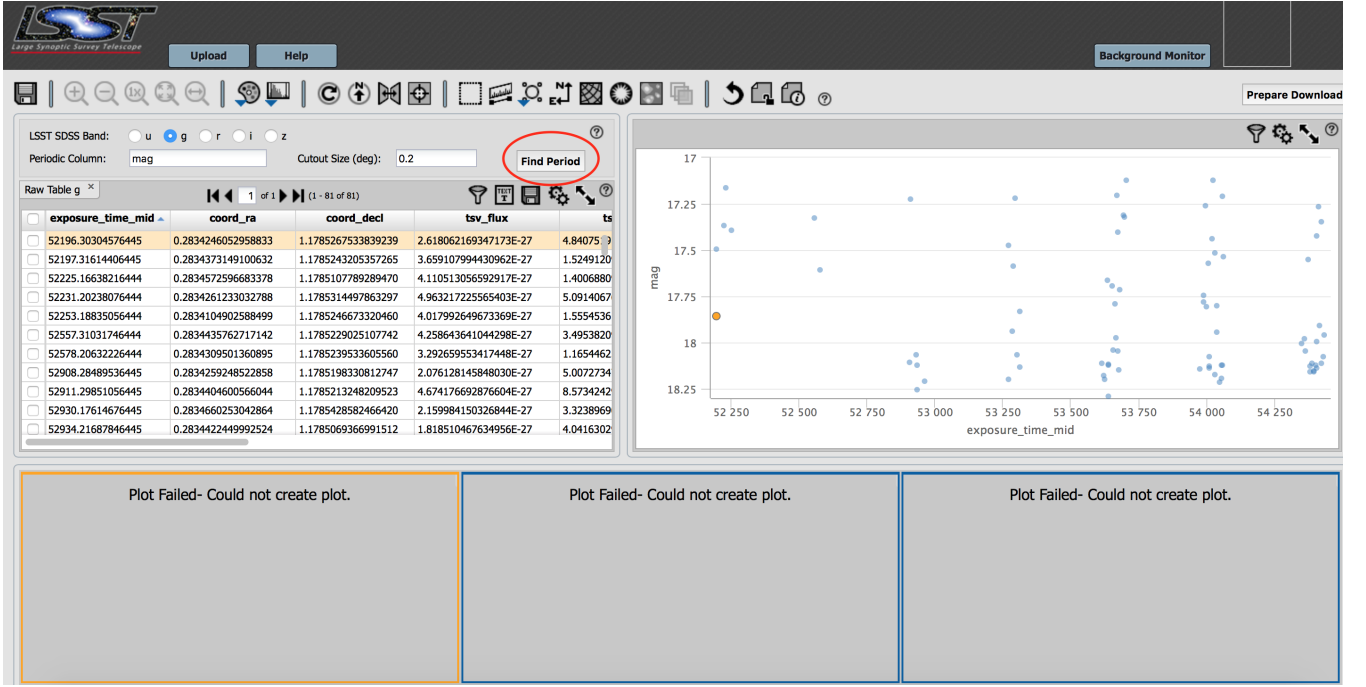


Figure 5. Time Series view for an RR Lyrae star at ra, dec = 0.283437, 1.178522 . Note that initially on the bottom there are three empty panels. The radio buttons in the upper left corner allow intuitively to select SDSS filter for lightcurve periodogram calculation (multi-band periodogram as in (VanderPlas & Ivezić 2015).) We select 'Find Period', marked with red oval, to calculate Lomb-Scargle periodogram for that band (this takes the user to Fig 6)

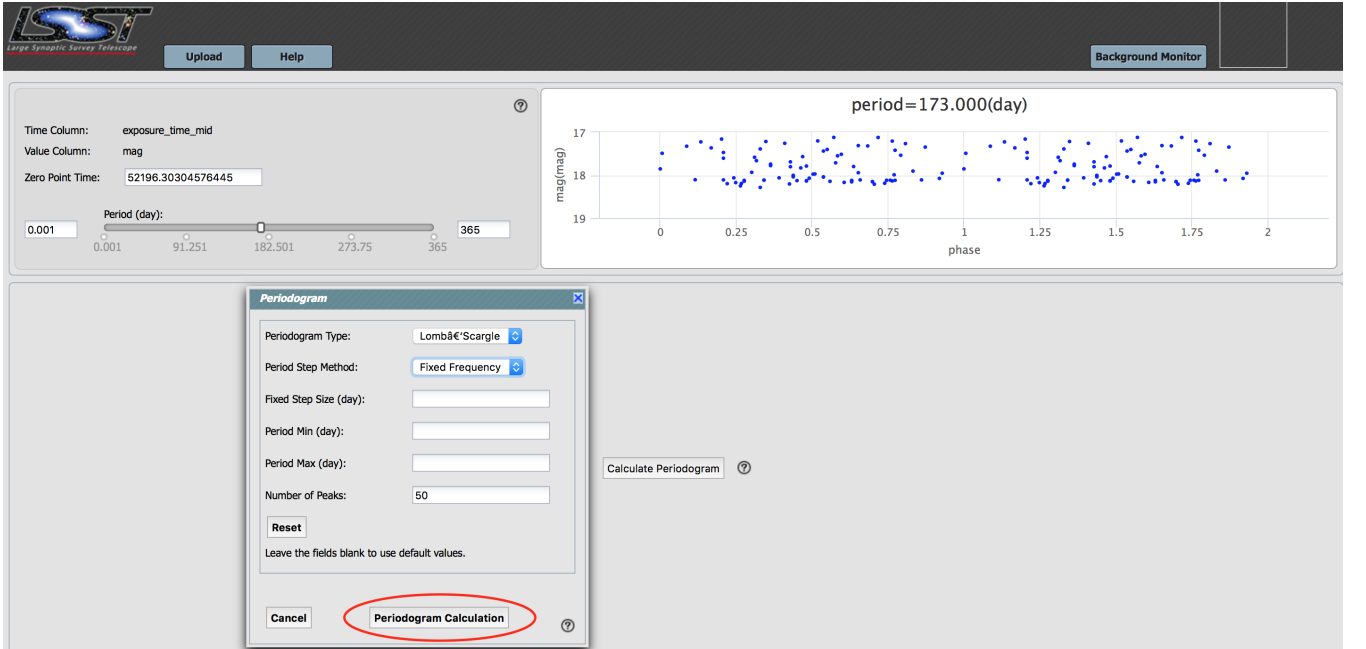


Figure 6. Calculating Lomb-Scargle periodogram for RR Lyrae at ra, dec = 0.283437, 1.178522 . The slider in the upper left corner allows to fold the lightcurve on the chosen period. Clicking 'Calculate Periodogram' button opens the dialog window 'Periodogram'. Currently it contains only the Lomb-Scargle as Periodogram Type. Period Step Method include Fixed Frequency or Fixed Period, similar to the NASA Exoplanet Periodogram Tool (Fig. 7). For this particular RR Lyrae we guide our choice of minimum and maximum periods (0.229, 0.998 days) by an educated guess, given that RR Lyrae in (Sesar et al. 2010) sample have periods between 0.254 and 0.907 days. Clicking on 'Periodogram Calculation' proceeds with evaluating Lomb-Scargle periodogram with chosen Period Step Method

NASA EXOPLANET ARCHIVE
NASA EXOPLANET SCIENCE INSTITUTE

Home About Us Data Tools Support Login

Periodogram Inputs Edit Input Table Plot Input Results

Periodogram Inputs

Input File Options	Algorithm and Period Settings
<p>Upload Data File:</p> <p>Choose File 13350_g.txt</p> <p>Upload</p> <p>Current Periodogram Data File:</p> <p>Name: 13350_g.txt</p> <p>Source: user uploaded file</p> <p>Edit Input Table</p> <p>Select Column Names:</p> <p>Time Column: col1</p> <p>Data Column: col2</p> <p>Plot Time vs. Data Columns</p> <p>Input File Information:</p> <p>Points used: 58 of 58</p> <p>Time range: 51075.302311 to 54412.235925</p> <p>Data range: 17.113 to 18.242</p>	<p>Select Algorithm:</p> <p>Algorithm: Lomb-Scargle</p> <p>Period Range:</p> <p>Minimum Period: 0.228731</p> <p>Maximum Period: 0.998246</p> <p>Period Step Method:</p> <p>Select Method: Fixed Frequency</p> <p>Fixed Step Size: 0.0001226</p>

Default(s) calculated successfully.

Calculate Periodogram Start New Session

Calculation Name: 13350_g.txt

Estimated processing time: 15 seconds

Figure 7. The same object as Fig. 9, and Fig 15, using the SDSS data from (Sesar et al. 2010). The highest significance frequency peak (power 21.58) corresponds to a period of 0.35365194 days. Only the second in significance peak (power 20.62) corresponds to the 'true' period of 0.547969 (Sesar et al. 2010). Note the bottom-left corner : the calculation took 15 secs for one lightcurve (compare to few milliseconds of AstroML code naive single-sinusoid approach that gave the same result for this particular object)

Fig. 10). We find that for about half of the SDSS lightcurves from Sesar et al. (2010), the Lomb-Scargle periodogram fitting single-term Fourier Series (LS) is sufficient to find the 'true' period, (see Fig. 12). We illustrate examples of RR Lyrae falling into each group : where with the naive LS we find the same period (Fig. 9), a smaller period (Fig. 10), , or a bigger period (Fig. 11) than the ground truth. For the same objects we also show PDAC lightcurves for which we also computed LS periodogram - see Figs. 15, 16 and 17, respectively.

Using the RA,Dec for the RR Lyrae from Sesar et al. (2010) we positionally query the PDAC RunDeepForcedSource database (Cone Search), to find objects within 2 arcsec radius. As shown on Fig. 13, not all objects have a match. For the 343 stars with a PDAC match, we obtain calibrated g-magnitude lightcurves querying the

RunDeepForcedSource and Science_Ccd_Exposure for the zero point magnitudes per exposure. For these PDAC lightcurves we also calculate Lomb-Scargle periodogram and find the frequency with most-significant power.

6 RESULTS

7 CONCLUSIONS

ACKNOWLEDGEMENTS

Thank you !




Figure 8. Distribution of RR Lyrae periods for 483 objects in (Sesar et al. 2010). Note the bimodal distribution, reflecting two main RR Lyrae types : 309 R Rab (right) and 104 R Rc (left) (see also Fig.16 in (Sesar et al. 2010)).

Table 1. Comparison of RR Lyrae periods obtained with different methods. First, P(S) is the 'ground truth' - period resulting from detailed template fitting by (Sesar et al. 2010). Second, P(LS) is the period corresponding to the most prominent frequency in the Lomb-Scargle periodogram (LS) computed on the SDSS lightcurve for a given object pulled from online journal data in (Sesar et al. 2010). Third, P(EXO) uses the same SDSS data from (Sesar et al. 2010) in g-band, to find the best period with the NASA Exoplanet Archive Periodogram service. Fourth, P(PDAC) uses the data pulled from PDAC, for which we find the best period using LS periodogram (same method as P(LS)).

ID	P(S)	P(LS)	P(EXO)	P(PDAC)
4099	0.641754	0.280827	0.64175	0.280827
13350	0.547987	0.547161	0.35365	0.547969
470994	0.346794	0.531667	0.34679	0.531667

REFERENCES

- Ivezić, Ž., Connolly, A. J., VanderPlas, J. T., & Gray, A. 2014, Statistics, Data Mining, and Machine Learning in Astronomy
- Sesar, B., et al. 2010, ApJ, 708, 717
- Vanderplas, J., Connolly, A., Ivezić, Ž., & Gray, A. 2012, in Conference on Intelligent Data Understanding (CIDU), 47
- VanderPlas, J., & Ivezić, Ž. 2015, The Astrophysical Journal, 812, 18
- VanderPlas, J. T. 2017, ArXiv e-prints
- Zechmeister, M., & Kürster, M. 2009, A&A, 496, 577




Figure 9. An example of the *astroML* Lomb Scargle periodogram performance, calculated for RR Lyr ID=13350 in SDSS g band (following Table 2 in (Sesar et al. 2010)), using the SDSS data from (Sesar et al. 2010). It took 18.6 milliseconds on a laptop to calculate this periodogram. The upper left panel depicts the raw SDSS lightcurve data. The upper right panel shows the phased lightcurve constructed with a cited period of 0.547987 days (P_{true}). The lower left panel shows the Lomb Scargle periodogram on a uniform frequency grid (5000 bins), where the orange and magenta vertical lines mark the location of the highest periodogram peak, and the frequency based on the reported period ($\omega_{true} = 2\pi/P_{true}$). The lower right panel shows the phased lightcurve constructed with the Lomb-Scargle Periodogram period of 0.547161 days, corresponding to the highest peak, $P_{fit} = 2\pi/\omega_{fit}$. The horizontal red and green lines mark the 5% and 1% significance levels for the highest peak, as found from 500 bootstrap resamplings (See http://www.astroml.org/book_figures/chapter10/index.html). The same object, but pulling the data from PDAC, is shown on Fig. 15




Figure 10. A periodogram, raw and folded lightcurve using SDSS data from (Sesar et al. 2010) for RR Lyr ID =4099 (the same object with PDAC data is shown on Fig. 16). It is an example of a failure of naive single Lomb Scargle periodogram performance - the ratio of $\omega_{true}/\omega_{fit} = 0.437$. The 'true' period from (Sesar et al. 2010) is 0.641754 days, whereas the naive Lomb-Scargle periodogram approach yields the 'fit' period of 0.280827 days. Note that here ω_{fit} and ω_{true} significantly differ for this RR Lyr, and the 'true' frequency, backed-up by the full lightcurve fitting of (Sesar et al. 2010), appears as only one of insignificant periodogram peaks. Everything else as on Fig. 9.




Figure 11. Same as Fig. 10, using the SDSS data pulled from (Sesar et al. 2010), with $\omega_{true}/\omega_{fit} = 1.53$. Here RR Lyr ID=470994 has a cited period of 0.346794 days (P_{true}), whereas period derived from the Lomb-Scargle periodogram is 0.531667. It may be a good example of aliasing.




Figure 12. The distribution of the ratio of ω_{true} to ω_{fit} , where ω_{true} is inferred directly from the 'ground truth' - period cited in Table 2 of (Sesar et al. 2010). We take the same SDSS data from the paper (Table 1 in (Sesar et al. 2010)), and calculate the Lomb-Scargle single-term generalized periodogram. The frequency corresponding to the highest peak is ω_{fit} . Thus, wherever this ratio is approximately equal to 1, this means that the naive LS approach is able to recover the 'true' period. However, where the highest frequency peak is not the same as ω_{true} , the ratio will be smaller or bigger from 1. This may be caused by the inherent simplicity of the simple single-term Fourier Series fitting. Indeed, some RR Lyrae lightcurves may have shapes that are insufficiently described by a single sinusoid (as on Fig.10.18 in (Ivezić et al. 2014)).




Figure 14. Comparison of RR Lyr ID=4099 from (Sesar et al. 2010) (red crosses), and PDAC (blue crosses). The two lightcurves have different length : 59 vs 162 points, respectively.




Figure 13. Results of positional query against 483 RR Lyrae stars from (Sesar et al. 2010), using their RA, Dec. Blue dots are 343 stars that have a match in the PDAC S82 dataset within 2 arcsec, and red crosses are 140 stars that did not. Increasing the search radius to 3 arcsec does not alter this result.




Figure 15. The same object as Fig. 9, but using data downloaded using PDAC. Using PDAC data, the RR Lyr ID=13350 has a best-fit period of 0.547969 days, almost identical to true period of 0.547987 from (Sesar et al. 2010). Panels the same as on Fig. 10




Figure 16. The same object as Fig. 10, but using data downloaded using PDAC. Calculating a naive LS periodogram using PDAC data for RR Lyr ID=4099 we find the best-fit period (frequency with highest power) of 0.280827 days, almost identical to the period found using LS periodogram on the SDSS (Sesar et al. 2010) data of 0.280827 days. Both are discrepant with respect to the 'true' period of 0.641754 days from (Sesar et al. 2010). Panels the same as on Fig. 9




Figure 17. The same object as Fig. 11, but using data downloaded from PDAC. Calculating a naive LS periodogram using PDAC data for RR Lyr ID=470994 we find the best-fit period (frequency with highest power) of 0.531667 days, almost twice as high as the 'true' period of 0.346794 days from (Sesar et al. 2010). For this star we get an identical period if we use LS periodogram on SDSS data from (Sesar et al. 2010) as opposed to PDAC. Panels the same as on Fig. 9

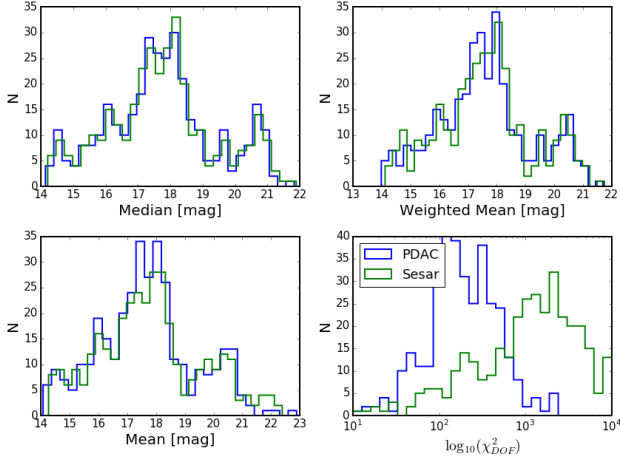


Figure 18. Comparison of the original (Sesar et al. 2010) lightcurves (green) against data for the same objects pulled from PDAC (blue). For each of the 383 lightcurves in SDSS *g*-band, without any pre-processing or clipping, we calculated the median, weighted mean, mean, and χ^2_{DOF} .

APPENDIX A: SQL QUERIES

To obtain lightcurves for objects with given RA,DEC, we would first query the RunDeepSource table to find out which objectIds are within the search radius, and then query the RunDeepForcedSource table against these objectIds. The two can be combined, querying for all lightcurves within a certain search radius.

A1 Cone query around given RA, DEC

To query RunDeepSource table, for an object with RA = 0.935679, and DEC=1.115859, search radius 0.00055555 (all in degrees), saving the output of the query to catalogCone1.json file:

```
1 curl -o catalogCone1.json -d 'query=SELECT+objectId+FROM+RunDeepSource+WHERE+\
2         qserv_areaspec_circle(0.935679,1.115859,0.00055555);' \
3         http://lsst-qserv-dax01.ncsa.illinois.edu:5000/db/v0/tap/sync
```

I perform such queries for each of the 483 objects in [Sesar et al. \(2010\)](#), using search radius of 2 arcsec.

A2 ID query to return lightcurves

We find that between 1 to 6 unique detection ID's are found within a search radius of 2 arsec for the 483 RR Lyrae. Given the IDs corresponding to the sought RA, DEC, location, to obtain the lightcurve in calibrated magnitudes, we need to query RunDeepForcedSource Table to provide information about individual exposures (mjd time of observation: exposure_time_mid, flux in a given filter flux_psf, flux error flux_psf_err, filter code exposure_filter_id, with 0,1,2,3,4 corresponding to u,g,r,i,z, respectively). We query scienceCcdExposureId table to obtain the flux for photometric zero point with associated error (fluxMag0, fluxMag0Sigma) with exposure_id in RunDeepForcedSource corresponding to scienceCcdExposureId in Science_Ccd_Exposure. We perform a join between these two tables, and store the result as catalogForced1.json file. Assuming that IDs that we query against are 3562429887808969, 3562429887812663, the query to obtain g magnitudes would be (see <https://confluence.lsstcorp.org/display/DM/PDAC+sample+queries+and+test+cases#>) :

```
1 curl -o catalogForced1.json -d 'query=\
2 SELECT \
3     objectId, id, fsrc.exposure_id, fsrc.exposure_time_mid, exp.run, \
4     scisql_dnToAbMag(fsrc.flux_psf,exp.fluxMag0) AS g, \
5     scisql_dnToAbMagSigma(fsrc.flux_psf, fsrc.flux_psf_err, \
6         exp.fluxMag0, exp.fluxMag0Sigma) AS gErr \
7 FROM \
8     RunDeepForcedSource AS fsrc, \
9     Science_Ccd_Exposure AS exp \
10 WHERE \
11     exp.scienceCcdExposureId = fsrc.exposure_id \
12     AND fsrc.exposure_filter_id=1 \
13     AND objectId IN (3562429887808969, 3562429887812663)\
14 ORDER BY exposure_time_mid' http://lsst-qserv-dax01.ncsa.illinois.edu:5000/db/v0/tap/sync
```

A3 Cone query to return lightcurves

We can combine the two steps above by querying positionally against the RunDeepForcedSource table :

```
1 curl -o catalogForced2.json -d 'query=\
2 SELECT \
3     objectId, id, fsrc.exposure_id, fsrc.exposure_time_mid, exp.run, \
4     scisql_dnToAbMag(fsrc.flux_psf,exp.fluxMag0) AS g, \
5     scisql_dnToAbMagSigma(fsrc.flux_psf, fsrc.flux_psf_err, \
6         exp.fluxMag0, exp.fluxMag0Sigma) AS gErr \
7 FROM \
8     RunDeepForcedSource AS fsrc, \
9     Science_Ccd_Exposure AS exp \
10 WHERE \
11     qserv_areaspec_circle(0.935679,1.115859,0.00055555) \
12     AND exp.scienceCcdExposureId = fsrc.exposure_id \
13     AND fsrc.exposure_filter_id=1 \
14 ORDER BY exposure_time_mid' http://lsst-qserv-dax01.ncsa.illinois.edu:5000/db/v0/tap/sync
```

A4 Box Query

We performed box query to compare ra,dec , objectId of all sources from a small region:

```

1
2 curl -o catalogBox0.json -d 'query=\
3 SELECT \
4     * \
5 FROM \
6     RunDeepSource \
7 WHERE \
8     qserv_areaspec_box(325,0,325.1,0.1);' \
9     http://lsst-qserv-dax01.ncsa.illinois.edu:5000/db/v0/tap/sync

```

APPENDIX B: READING JSON FILES

We used Python to read JSON files into AstroPy tables:

```

1 import json
2 from astropy.table import Table
3
4 json_file = 'catalogForced1.json'
5 tab = json.loads(open(json_file).read())
6 meta = tab['result']['table']['metadata']['elements']
7 data = tab['result']['table']['data']
8 dataT = np.transpose(data)
9
10 # initialize a new AstroPy Table
11 t = Table()
12
13 # iterate over columns to add them to the Table ...
14 for k in range(len(dataT)):
15     t.add_column(Column(dataT[k], name=meta[k]['name']), index=0)

```

This paper has been typeset from a T_EX/L^AT_EX file prepared by the author.